

```
In [1]: import tensorflow as tf
from tensorflow import keras
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
import pydot, graphviz
import numpy as np
import pandas as pd
```

Import datasets + simple exploration

```
In [2]: # import inbuilt fashion data from keras
fashion_mnist = keras.datasets.fashion_mnist
(X_train_full,y_train_full),(X_test,y_test)=fashion_mnist.load_data()
```


```
In [18]: # actual label names
label_names= {0:'T-shirt/top',
               1:'Trouser',
               2:'Pullover',
               3:'Dress',
               4:'Coat',
               5:'Sandal',
               6:'Shirt',
               7:'Sneaker',
               8:'Bag',
               9:'Ankle boot'}
```

```
In [4]: X_train_full.dtype
```

```
Out[4]: dtype('uint8')
```

```
In [19]: # show first image
plt.imshow(X_train_full[0])
```

```
Out[19]: (<matplotlib.image.AxesImage at 0x2b5f1ef9648>,,)
```



split into training and validation sets

- validation set is used for parameter tuning and model evaluation
- test set is used for final evaluation of performance

```
In [25]: # split using train_test_split function
X_train, X_valid, y_train, y_valid = train_test_split(X_train_n, y_train_full, test_size=0.2)
```


```
In [26]: # check dimensions
print('Training set Dimensions: \n\tX: {}\n\tY: {}'.format(X_train.shape, y_train.shape))
print('Validation set Dimensions: \n\tX: {}\n\tY: {}'.format(X_valid.shape, y_valid.shape))
print('Test set Dimensions: \n\tX: {}'.format(X_test.shape))
```

```
Training set Dimensions:
      X: (55000, 28, 28, 1)
      y: (55000,)
Validation set Dimensions:
      X: (5000, 28, 28, 1)
      y: (5000,)
Test set Dimensions:
      X: (10000, 28, 28, 1)
      y: (10000,)
```

Build Convolutional Neural Network

1. 3x3 filters are applied to give feature maps in the convolutional layer
 - these filters are trained to highlight on certain features
 - valid padding is applied, which ignores the input borders
2. Max pooling applied
 - max value in each grid is taken
 - this emphasizes on important features and reduces computational load too
3. input flattened to single array as input for ANN
4. Hidden layers
5. Output layers

```
In [27]: # make the results consistent by setting seed
np.random.seed(1)
tf.random.set_seed(1)
```



```
In [28]: # create NN object using sequential API
model=keras.models.Sequential()

# add convolutional layer
model.add(keras.layers.Conv2D(filters=32,
                               # i.e. filter size
                               kernel_size=(3,3),
                               # i.e. number of steps each filter takes
                               strides=1,
                               # this padding ignores input borders
                               padding='valid',
                               activation='relu',
                               # FOR COLORED IMAGES ,28x28x3 instead
                               input_shape=(28,28,1)))

# max pooling taking max value from 2x2 grid
model.add(keras.layers.MaxPooling2D((2,2)))

# flatten into a single array
model.add(keras.layers.Flatten())

# add a hidden layer of 300 neurons with relu activation
model.add(keras.layers.Dense(300,activation='relu'))
# add a hidden layer of 100 neurons with relu activation
model.add(keras.layers.Dense(100,activation='relu'))

# add output layer of 10 neurons(there are 10 classes) with softmax activation
model.add(keras.layers.Dense(10,activation='softmax'))
```

```
In [29]: model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_1 (MaxPooling2D)	(None, 13, 13, 32)	0
flatten_1 (Flatten)	(None, 5408)	0
dense_3 (Dense)	(None, 300)	1622700
dense_4 (Dense)	(None, 100)	30100
dense_5 (Dense)	(None, 10)	1010
Total params: 1,654,130		
Trainable params: 1,654,130		
Non-trainable params: 0		

- Param:
 - refers to the number of parameters the model needs to tune

```
In [31]: # setting parameters that are problem specific
model.compile(loss='sparse_categorical_crossentropy',
              # sgd refers to stochastic gradient descent
              optimizer='sgd',
              metrics=['accuracy'])
```

```
In [41]: # Load saved model
model = keras.models.load_model('CNNmodel.h5')
```

```
In [42]: # using .fit method consecutively actually continues training the model where it was left
model_history = model.fit(X_train,y_train,
                           validation_data=(X_valid,y_valid),
                           epochs=30)
```

Train on 55000 samples, validate on 5000 samples

Epoch 1/30

55000/55000 [=====] - 30s 551us/sample - loss: 0.1335 - accuracy: 0.9528 - val_loss: 0.2452 - val_accuracy: 0.9126

Epoch 2/30

55000/55000 [=====] - 44s 808us/sample - loss: 0.1274 - accuracy: 0.9545 - val_loss: 0.2488 - val_accuracy: 0.9128

Epoch 3/30

55000/55000 [=====] - 42s 772us/sample - loss: 0.1232 - accuracy: 0.9559 - val_loss: 0.2314 - val_accuracy: 0.9176

Epoch 4/30

55000/55000 [=====] - 43s 774us/sample - loss: 0.1178 - accuracy: 0.9579 - val_loss: 0.2513 - val_accuracy: 0.9088

Epoch 5/30

55000/55000 [=====] - 44s 801us/sample - loss: 0.1124 - accuracy: 0.9610 - val_loss: 0.2370 - val_accuracy: 0.9156

Epoch 6/30

55000/55000 [=====] - 43s 782us/sample - loss: 0.1072 - accuracy: 0.9620 - val_loss: 0.2353 - val_accuracy: 0.9176


```
In [33]: # show model parameters
model_history.params
```

```
Out[33]: {'batch_size': 32,
          'epochs': 30,
          'steps': 1719,
          'samples': 55000,
          'verbose': 0,
          'do_validation': True,
          'metrics': ['loss', 'accuracy', 'val_loss', 'val_accuracy']}
```

```
In [34]: # show the training progression
model_history.history
```

```
Out[34]: {'loss': [0.7232647918701172,
                  0.4789659946875139,
                  0.4240220858747309,
                  0.38848866054795006,
                  0.36352404016581447,
                  0.3416003832513636,
                  0.3246729166767814,
                  0.30834569490606134,
                  0.2958024963920767,
                  0.2845826054659757,
                  0.2724606364792044,
                  0.26324330543821506,
                  0.2536492290626873,
                  0.24500095049684698,
                  0.23684471851587297,
                  0.22951359150409698,
                  0.22168787875717336,
                  0.2143544225161726,
                  0.2067171369996938,
                  0.20066402100573016]
```

```
In [43]: # visualise the training progression
# after 60 epochs the model is near 100% accuracy
pd.DataFrame(model_history.history).plot(figsize=(8,5))
plt.grid(True)
plt.gca().set_ylim(0,1)
plt.show()
```



Evaluating performance

```
In [36]: # 0.318 Loss and accuracy of 0.889
model.evaluate(X_test,y_test)
```

10000/10000 [=====] - 1s 140us/sample - loss: 0.2543 - accuracy: 0.9121

```
Out[36]: [0.2543268059760332, 0.9121]
```

```
In [37]: # making single probability predictions
y_proba = model.predict(X_test[:3])
y_proba.round(2)
```

```
Out[37]: array([[0., 0., 0., 0., 0., 0., 0., 0., 0., 1.],
                [0., 0., 1., 0., 0., 0., 0., 0., 0., 0.],
                [0., 1., 0., 0., 0., 0., 0., 0., 0., 0.]], dtype=float32)
```


```
In [38]: # or we could make class predictions
y_pred = model.predict_classes(X_test[:3])
y_pred
```

```
Out[38]: array([9, 2, 1], dtype=int64)
```

```
In [41]: print(label_names[y_pred[0]])  
plt.imshow(X_test[0])
```

Ankle boot


```
Out[41]: <matplotlib.image.AxesImage at 0x20375271f88>
```



```
In [42]: print(label_names[y_pred[1]])  
plt.imshow(X_test[1])
```

Pullover


```
Out[42]: <matplotlib.image.AxesImage at 0x203745d4408>
```




```
In [43]: print(label_names[y_pred[2]])  
plt.imshow(X_test[2])
```

Trouser

```
Out[43]: <matplotlib.image.AxesImage at 0x20375668048>
```



Saving and restoring

```
In [39]: # save model  
model.save('CNNmodel.h5')
```

```
In [40]: # delete model  
del model  
keras.backend.clear_session()
```

```
In [41]: # Load model  
model = keras.models.load_model('CNNmodel.h5')
```

```
In [57]: # Loaded model is the same as the model trained in previous session  
model.evaluate(X_test,y_test)
```

```
10000/10000 [=====] - 1s 109us/sample - loss: 0.3183 - accuracy: 0.8889
```

```
Out[57]: [0.3182890306711197, 0.8889]
```

